

FORCE MEMORY TRANSFORMERS: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study transformer robot policies. The motivating bet is: Give transformers persistent force-memory channels that survive visual scene changes.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For transformer robot policies, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2026; 2025a; 2022a; 2023; 2025b; 2022b). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.512	0.488
observed_only_wam	hidden_mode_shift	0.193	0.807
augmentation_as_data_wam	nominal	0.665	0.335
augmentation_as_data_wam	hidden_mode_shift	0.263	0.738
scalar_uncertainty_wam	nominal	0.623	0.377
scalar_uncertainty_wam	hidden_mode_shift	0.380	0.620
proposed_branch_mechanism	nominal	0.782	0.218
proposed_branch_mechanism	hidden_mode_shift	0.590	0.410

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Force Memory Transformers argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Semantic-geometric visual place recognition: a new perspective for reconciling opposing views, 2022a. <https://doi.org/10.1177/0278364919839761>.
- Masked vision-language transformers for scene text recognition, 2022b. <http://arxiv.org/abs/2211.04785v1>.
- Surgical-vqla:transformer with gated vision-language embedding for visual question localized-answering in robotic surgery, 2023. <https://doi.org/10.1109/icra48891.2023.10160403>.
- Scene-llm: Extending language model for 3d visual reasoning, 2025a. <https://doi.org/10.1109/wacv61041.2025.00220>.
- Evaluating multimodal vision-language model prompting strategies for visual question answering in road scene understanding, 2025b. <https://doi.org/10.1109/wacvw65960.2025.00115>.
- Zerorel: Multimodal transformer-guided zero-shot relationship retrieval for generalized scene graph generation, 2026. <https://doi.org/10.1007/s11760-026-05439-7>.